



Celery

A Distributed Task Queue

Idan Gazit

PyWeb-IL 8 / 29th September 2009

What is Celery?



Celery is a...

Distributed
Asynchronous
Task Queue
For Django



Celery is a...

Distributed
Asynchronous
Task Queue
For Django



Celery is a...

Distributed
Asynchronous
Task Queue
For Django



Celery is a...

Distributed
Asynchronous
Task Queue
For Django



Celery is a...

Distributed
Asynchronous
Task Queue
For Django



What can I use it for?



Potential Uses


- » **Anything that needs to run asynchronously, e.g. outside of the request-response cycle.**
- » Background computation of 'expensive queries' (ex. denormalized counts)
- » Interactions with external API's (ex. Twitter)
- » Periodic tasks (instead of cron & scripts)
- » Long-running actions with results displayed via AJAX.



How does it work?




Celery Architecture



Celery Architecture

user




submit:
tasks
task sets
periodic tasks
retryable tasks



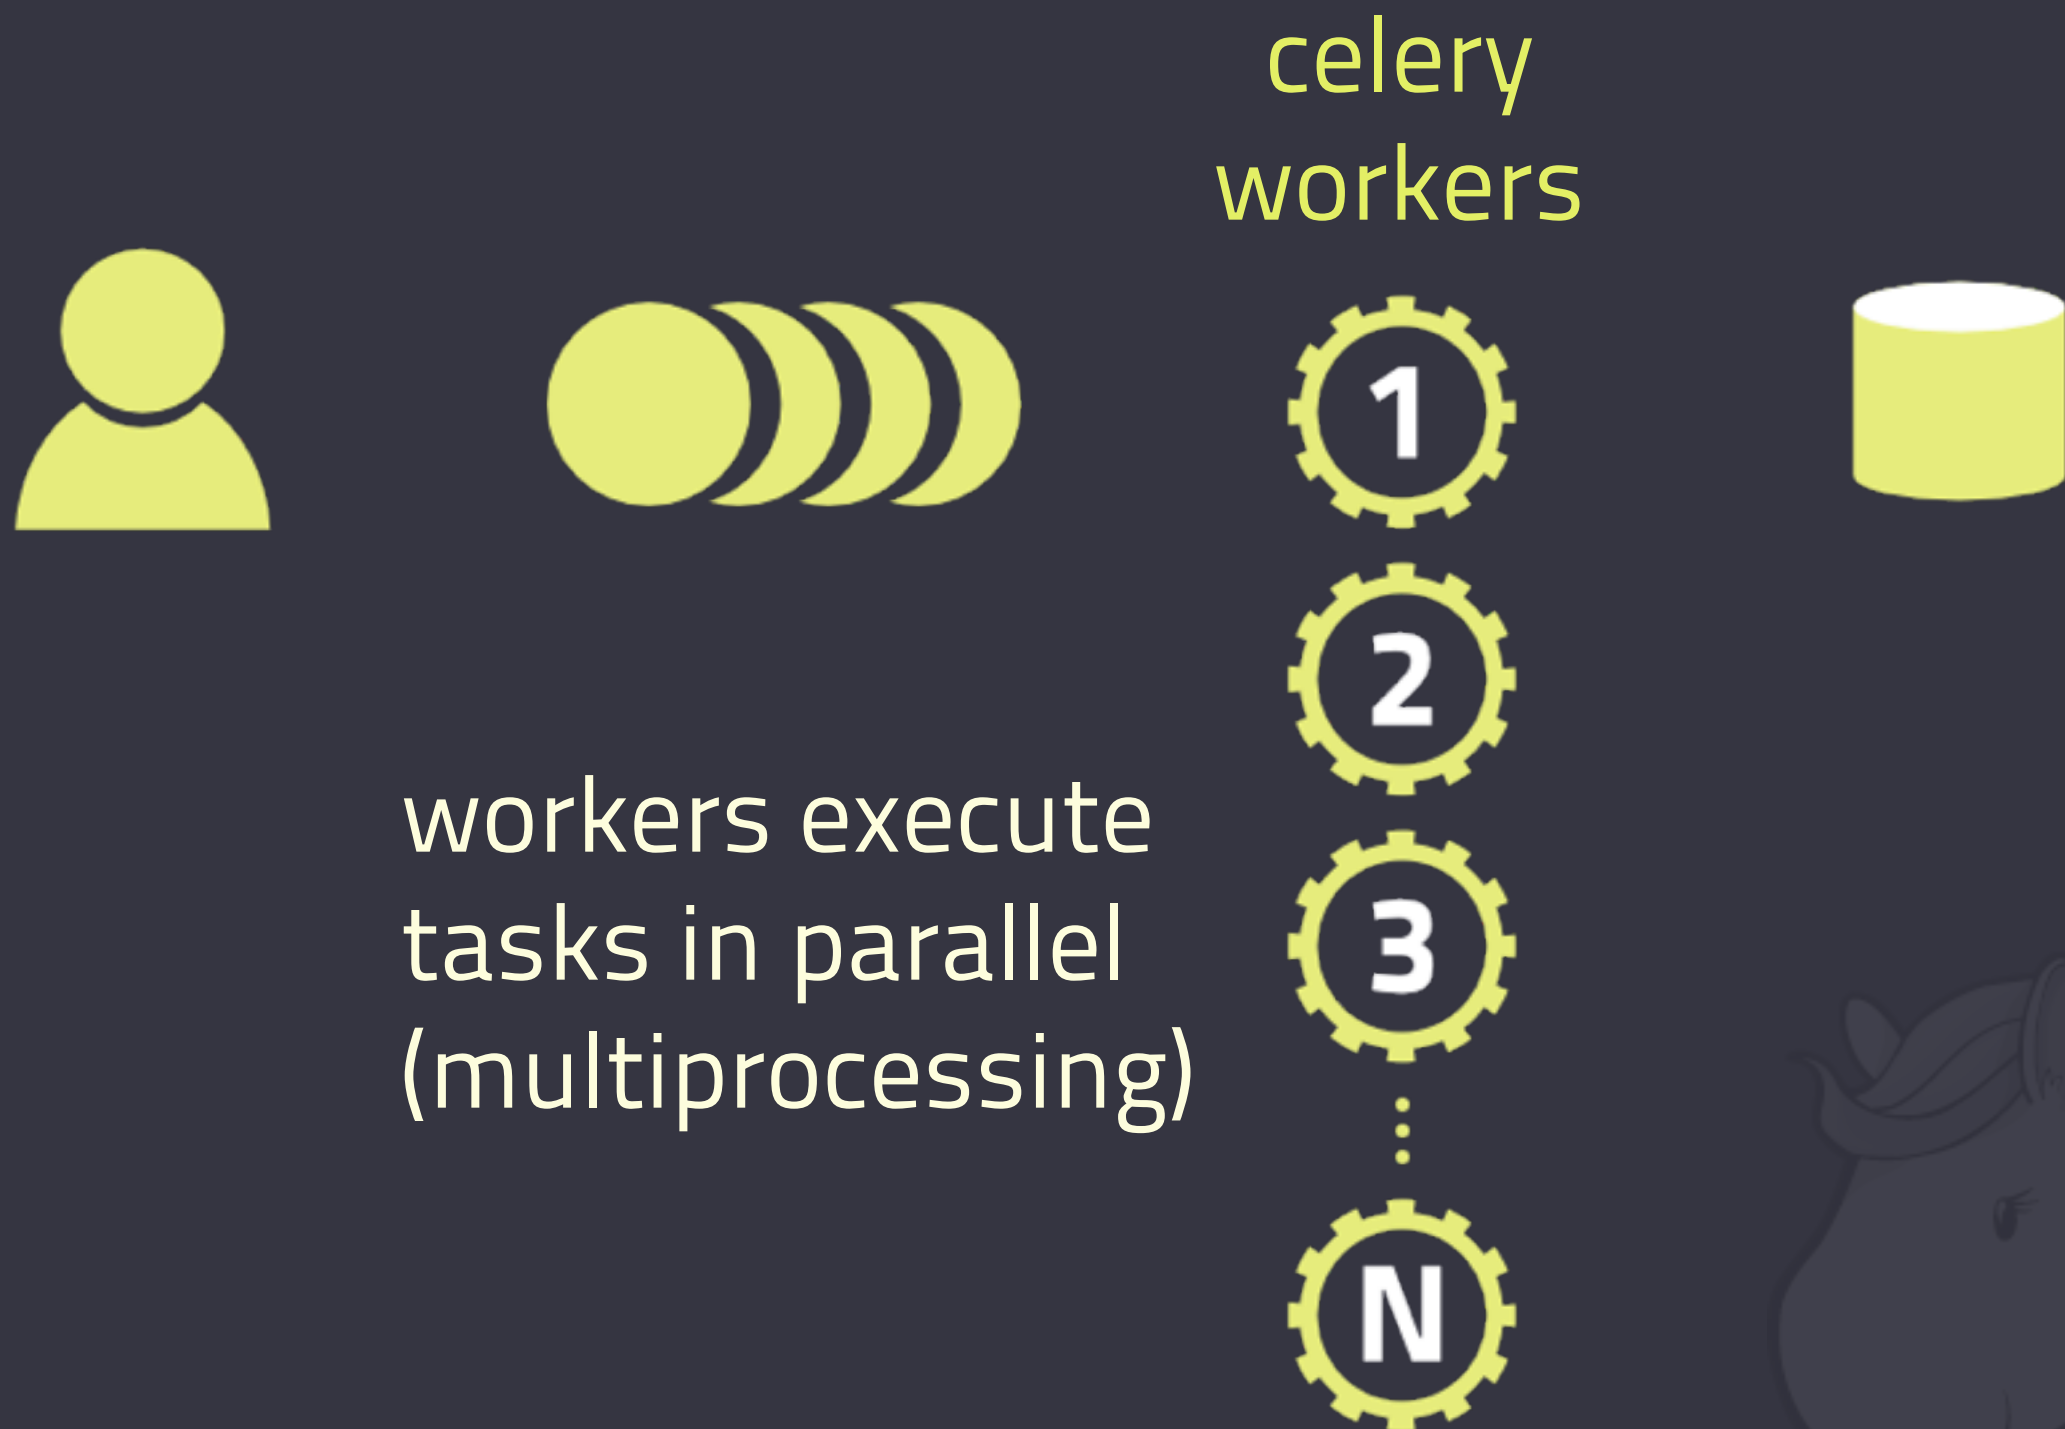
⋮



Celery Architecture



Celery Architecture



Celery Architecture



task result (tombstone)
is written to task store:

- ▶ RDBMS
- ▶ memcached
- ▶ Tokyo Tyrant
- ▶ MongoDB
- ▶ AMQP (new in 0.8)

celery
workers




...



task result
store



Celery Architecture



Celery Architecture

Celery

uses...

Carrot

to talk to...

**AMQP
Broker**



Celery Architecture

Celery

pip install celery

Carrot

(dependency of celery)

**AMQP
Broker**



Celery Architecture

Celery

`pip install celery`

Carrot

(dependency of celery)

RabbitMQ

<http://www.rabbitmq.com>




AMQP is... Complex.



AMQP is Complex

- » VHost
- » Exchanges
 - » Direct
 - » Fanout
 - » Topic
- » Routing Keys
- » Bindings
- » Queues
 - » Durable
 - » Temporary
 - » Auto-Delete



bit.ly/amqp_intro



I Can Haz Celery?




Adding Celery

1. get & install an AMQP broker
(pay attention to vhosts, permissions)
2. add Celery to INSTALLED_APPS
3. add a few settings:
AMQP_SERVER = "localhost"
AMQP_PORT = 5672
AMQP_USER = "myuser"
AMQP_PASSWORD = "mypassword"
AMQP_VHOST = "myvhost"
4. manage.py syncdb



Celery Workers

- » Run at least 1 celery worker server
- » `manage.py celeryd`
(`--detach` for production)
- » Can be on different machines
- » Celery guarantees that tasks are only executed once



Tasks



Tasks

- » Define tasks in your app
- » `app_name/tasks.py`
- » `register & autodiscovery`
(like `admin.py`)




Task

```
from celery.task import Task
from celery.registry import tasks

class FetchUserInfoTask(Task):
    def run(self, screen_name, **kwargs):
        logger = self.get_logger(**kwargs)
        try:
            user = twitter.users.show(id=screen_name)
            logger.debug("Successfully fetched {0}".format(screen_name))
        except TwitterError:
            logger.error("Unable to fetch {0}: {1}".format(
                screen_name, TwitterError))
            raise

        return user

tasks.register(FetchUserInfoTask)
```



Run It!

```
>>> from myapp.tasks import FetchUserInfoTask  
>>> result = FetchUserInfoTask.delay('idangazit')
```



Task Result

- » **result.ready()**
true if task has finished
- » **result.result**
the return value of the task or exception instance if the task failed
- » **result.get()**
blocks until the task is complete then returns result or exception
- » **result.successful()**
returns True/False of task success



Why even check results?



Chained Tasks

```
from celery.task import Task
from celery.registry import tasks

class FetchUserInfoTask(Task):
    def run(self, screen_name, **kwargs):
        logger = self.get_logger(**kwargs)
        try:
            user = twitter.users.show(id=screen_name)
            logger.debug("Successfully fetched {0}".format(screen_name))
        except TwitterError:
            logger.error("Unable to fetch {0}: {1}".format(
                screen_name, TwitterError))
            raise
        else:
            ProcessUserTask.delay(user)

        return user

tasks.register(FetchUserInfoTask)
```



Task Retries



Task Retries


```
from celery.task import Task
from celery.registry import tasks

class FetchUserInfoTask(Task):
    default_retry_delay = 5 * 60 # retry in 5 minutes
    max_retries = 5

    def run(self, screen_name, **kwargs):
        logger = self.get_logger(**kwargs)
        try:
            user = twitter.users.show(id=screen_name)
            logger.debug("Successfully fetched {0}".format(screen_name))
        except TwitterError, exc:
            self.retry(args=[screen_name,], kwargs=**kwargs, exc=exc)
        else:
            ProcessUserTask.delay(user)

        return user

tasks.register(FetchUserInfoTask)
```



Periodic Tasks



Periodic Tasks

```
from celery.task import PeriodicTask
from celery.registry import tasks
from datetime import timedelta

class FetchMentionsTask(Task):
    run_every = timedelta(seconds=60)

    def run(self, **kwargs):
        logger = self.get_logger(**kwargs)
        mentions = twitter.statuses.mentions()
        for m in mentions:
            ProcessMentionTask.delay(m)

        return len(mentions)

tasks.register(FetchMentionsTask)
```



Task Sets



Task Sets

```
>>> from myapp.tasks import FetchUserInfoTask
>>> from celery.task import TaskSet
>>> ts = TaskSet(FetchUserInfoTask, args=(
    ['ahikman'], {},
    ['idangazit'], {},
    ['daonb'], {},
    ['dibau_naum_h'], {}))
>>> ts_result = ts.run()
>>> list_of_return_values = ts.join()
```



MOAR SELRY!



Celery.Views



Celery.Views

- » Celery ships with some django views for launching / getting the status of tasks.
- » JSON views perfect for use in your AJAX (err, AHA) calls.
- » `celery.views.apply(request, task_name, *args)`
- » `celery.views.is_task_done(request, task_id)`
- » `celery.views.task_status(request, task_id)`



Routable Tasks



Routeable Tasks

>> "I want tasks of type X to only execute on this specific server"

>> Some extra settings in settings.py:

```
CELERY_AMQP_EXCHANGE = "tasks"  
CELERY_AMQP_PUBLISHER_ROUTING_KEY = "task.regular"  
CELERY_AMQP_EXCHANGE_TYPE = "topic"  
CELERY_AMQP_CONSUMER_QUEUE = "foo_tasks"  
CELERY_AMQP_CONSUMER_ROUTING_KEY = "foo.#" 
```

>> set the task's routing key:

```
class MyRouteableTask(Task):  
    routing_key = 'foo.bars'
```



like django, it's just python.




Support

#celery on freenode

<http://groups.google.com/group/celery-users/>

AskSol (the author) is friendly & helpful



Fin.

@idangazit
idan@pixane.com

